

Graph Models of Named Entity Types for Interpretable Named Entity Disambiguation

Andrew Larimer

andrewlarimer@berkeley.edu

Daniel Rasband

danrasband@berkeley.edu

Abstract

Like many neural networks, the decisions made by most Named Entity Recognition systems are difficult to interpret or adjust. We prototype a new Named Entity Recognition architecture that is very transparent and could be manually adjusted. Additionally, it is syntactically aware and focuses only on words with a direct dependency relationship to the entity in question, filtering out the noise of nearby words. Finally, our system builds a clear and cumulative ‘mental model’ of each named entity type, which may be extensible as a structure for persistent knowledge across long texts or for longterm knowledge-building.

1 Introduction

Named Entity Recognition systems have achieved impressive performance with state-of-the-art models resulting in $\sim 90\%$ classification accuracy on popular benchmarks like the OntoNotes 5.0 corpus and the CoNLL-2003 shared task data.

In order to achieve these results, however, the leading neural systems rely on matrices of numerical weights, often distributed across neural layers. This makes their rationale for arriving at a given classification uninterpretable by humans. While word embeddings, another mathematical construct of Natural Language Processing, can at least be explored in relation to one another by projecting them into 2D or 3D space, current NER systems don’t provide us an easy way to compare the internal representations of different entity types. In industries and applications that require decisions to be auditable, for example in the processing of loan applications that fall under the jurisdiction of the Fair Lending Act, the inability to peer into the workings of a model and clearly articulate why it

Model	Corpus	Type	Accuracy
Lample LSTM-CRF	CoNLL	neural	90.94
Lample Stack-LSTM	CoNLL	neural	90.33
Chiu and Nichols	OntoNotes	neural	86.19
Ratinov and Roth	OntoNotes	linear	83.45
Our Model	OntoNotes	graph	25.62
Baseline	OntoNotes	baseline	20.15

Table 1: Our model’s accuracy currently lags far behind other architectures, but surpasses a baseline of always choosing the most common entity type (ORG). The corpuses, more specifically, are CoNLL-2003 and OntoNotes 5.0.

is outputting one classification or another may not be acceptable.

We propose a novel graph-based architecture that allows for transparency in how a candidate entity is classified. While our prototype system achieves only a very modest accuracy of 25.62% on the OntoNotes 5.0 corpus, we are confident that it can be further enhanced through the incorporation of additional input features and refinements. The basic validity of the approach is demonstrated by our model’s statistically significant performance (p -value: $2.2e-16$) over a baseline of always choosing the most common Named Entity Type (‘ORG’ in the OntoNotes 5.0 corpus) which would result in an accuracy of 20.15%.

Our approach begins by training what we refer to as a Named Entity Template (NET) Graph. These graphs are intended to represent a ‘mental model’ or Platonic ideal of how this entity type interacts with the world (or at least how it is described as interacting with the world within text).

We achieve our results by first constructing a graph over a labeled training set that serves as a ‘mental model’ or ‘Platonic ideal’ of what a given

Named Entity Type is. Concretely, Figure 1 shows an example of a dramatically pruned version of our representation of the 'PERSON' entity type. The graphs are built to be grammatically aware, constructing a subgraph for each syntactic role the entity plays within the corpus. This allows for us to distinguish between what a given entity does in comparison to what is done to it.

In Figure 1, we see that when a 'PERSON' is the subject of a verb with an active tense, they might "concede", "dare", or "sign" something. When they are a direct object, someone might be "urging" them, "bringing" them somewhere, or "handing" them something. If used as a possession modifier, a person might have a "widow", a "son", or she might have "words".

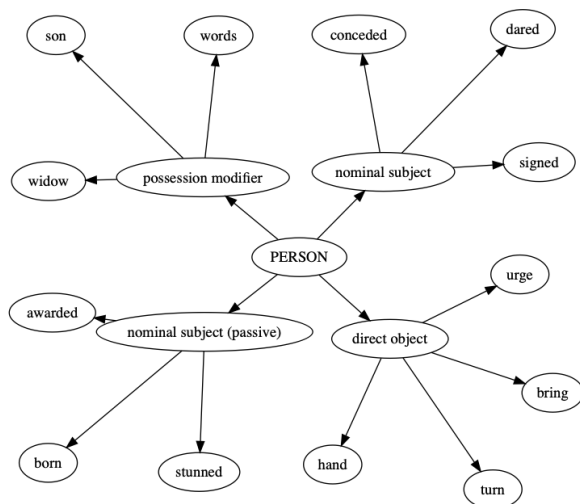


Figure 1: A simplified version of a Named Entity Type graph.

The example in Figure 1 shows a NET graph constructed from twelve occurrences of entities labeled as 'PERSON' in the OntoNotes corpus. We use a dependency parser to parse each sentence in which a labeled entity occurs, and we create a leaf node from the word identified by the dependency parse as the head of our labeled entity. The dependency relationship between them, which is typically labeled on the arc in a visualized dependency parse, becomes another node in between the root 'PERSON' node and its head in that occurrence.

When conducting inference, we currently rely on an existing Named Entity Recognition system to supply candidates, and our system takes over to classify the named entity type of the candidate.

For each candidate entity, we construct a graph in a similar manner to how we construct the

Named Entity Template graph. The primary difference is that while we have many labeled examples of a 'PERSON' in the OntoNotes corpus, when encountering a candidate entity, we only have its one concrete occurrence. To supplement its graph, we use a coreference system to find other mentions in the text that represent the same entity, and add the relationships from each of those mentions to our graph.

We note here that while Figure 1 shows only one generation of relationships in the direction of the dependency head, from a labeled entity to its direct head node according to a dependency parse, our NET graphs and candidate graphs actually traverse the dependency parse two generations up the tree (to the entity's head and that head's head) as well as two generations down the tree (to the entity's children, such as any adjectives modifying it, as well as the children of those children). This allows us to take in much more signal from the text that is still directly related to the entity we are considering.

Once a candidate graph has been construction, we compare it against each of our Named Entity Template graphs through a graph-similarity algorithm we developed. For our prediction we choose the label with the greatest similarity. In the case of a tie in maximum similarity scores, we choose the entity that occurs most frequently in the corpus from amidst the top scores.

We describe our process for building the graphs and computing their similarity further in Section 3.

We think the virtue of this approach lies primarily in the transparency of the model. If desired, the system could easily be extended to include a reference point in the text for each node, for example, to point a reader to exactly which mentions in the corpus contributed to a model's interpretation.

Additionally, like any model, the more data it encounters the more robust each NET graph will become, and we believe that the interpretability and generalizability of these graphs in describing the behavior and attributes of various entities might lead to them being utilized in other systems.

2 Background and Related Work

While we ultimately decided to build a new architecture, we were informed and inspired by previous work directly related to NER as well as work from other subfields. In this section, we provide

a brief survey of a few research papers and techniques and how they influenced our thinking.

2.1 Ontology-based NER

As our specific approach might be described as existing somewhere between Named Entity Recognition and Knowledge Representation, we considered approaches from both subfields.

We found the work of [Suciu and Groza \(2014\)](#) interesting in its Ontology-based approach for logically adding labels to characters in folktales by their role within the story. And how they relate to the other characters. For example, if a boy has a parent that is a King or a Queen, then that boy is a Prince:

$$Prince \equiv Boy \cap (\exists hasParent.King \cup \exists hasParent.Queen) \quad (1)$$

We found that an intriguing concept, and wondered if we could build a system that could classify entities not just by known facts such as having a parent who is a King or Queen, but by the way they act, are acted upon, or are described.

2.2 Transition-based NER

Within the subfield of NER, [Lample et al. \(2016\)](#) demonstrated interesting results in transition-based Named Entity Recognition. Using a similar model to a transition-based dependency parser, which processes text from left to right by considering one word at a time from a buffer of all the words in a sentence.

For each word it encounters, a traditional arc-standard transition-based system ([Nivre, 2004](#)) classifies the appropriate action to take from a small vocabulary of possible actions. One action in a transition-based system is a 'SHIFT' to move a word from the buffer onto a stack, which serves as both a storage area for words on which we are not yet ready to act, as well as a staging area for multi-word phrases. Another action in a typical dependency parsing transition-based system is 'REDUCE', which in arc-eager systems can come in a 'REDUCE-LEFT' and 'REDUCE-RIGHT' variety, and removes words from the buffer or stack depending on which directional variant is used, applying a dependency relation to the arc at the same time.

In this application, [Lample et al. \(2016\)](#) used a similar system enhanced by Stack-LSTMs([Dyer et al., 2015](#)) instead of simple stacks. An LSTM is a type of recurrent neural network that passes

two outputs into another LSTM cell that shares its weights. One of these outputs is more directly responsive to the immediate input at each word as it moves over a text, while the other output called the 'cell state' is designed to carry signal until the chain of LSTM cells decides to alter it, thus enabling a longer 'memory' of the passage of text that has been fed through the network. The Stack-LSTMs used in this case are a hybrid of a stack data structure and a recurrent LSTM neural network designed such that each stack can maintain an embedded awareness of its full contents. And here, instead of applying dependency parse tags, their weights are trained to apply NER type labels upon a REDUCE action.

Transition-based approaches were initially designed to generate trees, and while the weights for applying labels are opaque to human interpretability, we found this method intriguing and inspiring towards our eventual approach.

2.3 LSTMs and CRFs

While we were particularly interested in the transition-based approach, we would be remiss in failing to discuss another architecture from the same paper([Lample et al., 2016](#)) which had superior performance: a combination of bi-directional LSTMs and Conditional Random Fields.

In this model, the output of a pair of bi-directional LSTMs fed their output into a Conditional Random Field, which is a context-aware structure that is able to jointly model across the inputs from the sequence to capture the most likely NER tags across the sequence.

While their results of 90.46% on the CoNLL-2003 corpus are impressive, both these results and those of the previous approach suffer from the lack of interpretability we discussed earlier, and thus we chose an alternate approach.

3 Methods

We chose to use the OntoNotes 5.0 corpus for its sizable corpus of 143,709 sentences with labeled named entities. We divided the corpus into training, dev, and test sets with a ratio of 70 / 20 / 10.

3.1 Creating our Named Entity Type Graphs

For each occurrence of a labeled named entity in our training set, we do the following:

1. If this occurrence of a named entity consists of more than one word (ex. 'Nobel Peace Prize'),

we assume the head of the phrase to be the last word of the named entity phrase.

2. We load the graph for the Named Entity Type represented by that occurrence ('PERSON', 'ORGANIZATION', 'WORK_OF_ART', etc.). Using the example of 'Nobel Peace Prize', we would load the graph for its entity type: 'WORK_OF_ART'.

3. We look at the dependency parse of the sentence in which the entity occurrence appears, and add the entity's head token to a list, and its children token to a list. The steps below repeat for each of the words added to these lists.

4. We look at the relationship identified by the dependency parse between the root node of the graph and the head or child we are considering. If the Named Entity Type graph already has a node corresponding to that relationship, we load it as our "intermediary node" and add a one to its occurrence count. If it does not yet exist, we create it and assign it a count score of one. In the 'Nobel Peace Prize' example, we see that 'Prize' is a 'direct object' with 'awarded' as its head, so we find the 'direct object' child node of the 'WORK_OF_ART' root node. If there isn't already a 'direct object' node, we create it, and assign it a score count of one. If it is already there, we increase a score count by one for the ability to explore different similarity functions later on.

We chose to implement the dependency relationships as intermediary nodes rather than labels on our edges for two reasons: First, it saves us time during inference by allowing us to use graph library methods to directly access the intermediary node that represents how a given candidate entity is being used in a sentence. Additionally, it generates cleaner, more interpretable graphs, with natural clustering of the head words resulting from a particular syntactic usage of the entity type.

5. Under the intermediary node, we now add a leaf node for our named entity's head in the dependency parse. In the 'Nobel Peace Prize' example, the head of 'Prize' when it was used as a direct object is 'awarded,' so we create a node for 'awarded' as a child node of 'direct object' and give it a score of one, unless it already exists, in which case we increase its score by one.

6. In the special case of arriving at a preposition, we take an extra leap up the dependency tree to find that preposition's head, as we found more signal present in the heads of the preposi-

tions rather than in the prepositions like "to" or "of" themselves.

7. If we were processing a head node, we add its head to our head list, and if we were processing a child node, we add its children to our child list, and we repeat the process for one more generation, with this node we have arrived at acting as the new root node of our next generation.

3.2 Conducting Inference on a New Named Entity Candidate

We consider each Named Entity as identified by an existing NER system¹ For each named entity candidate, we do the following:

1. We construct a candidate graph in a similar manner to that described above for generating the Named Entity Type graphs. We use a coreference system to include multiple references to the entity in order to generate as robust a graph as possible to compare against our NET graphs.

2. We compare our candidate graph to each of the eighteen Named Entity Type graphs. We only consider the subgraphs branching from intermediary nodes (the syntactic role nodes from the dependency relations) that are contained in our candidate graph, which saves us a lot of processing time on each of the extensive NET graphs.

3. Traveling down the branches of the matching intermediary nodes, we calculate a score for each successor node (the words that represent the head or child of our original node). If that successor node is a leaf node, we calculate the score directly as described below. If it has further intermediary node successors (which would be further dependency relations that match between the candidate graph and the NET graph against which we are comparing it), we travel down those branches to their leaf nodes, and eventually assign this node a score of the sum of its successors' scores.

4. For each of our candidate graph's leaf nodes, we compare that word's similarity to the similarity of all the leaf nodes under the NET's corresponding intermediary node by a comparison of their word embeddings.² We take the highest-scoring similarity measurement as our score for this leaf node. Concretely, if we were considering

¹For candidate identification we used spaCy's (Honnibal and Montani, 2017) built-in NER module.

²This score is using spaCy's Token.similarity() method. This seems to be a cosine similarity of the words' embeddings, but it is not explicitly stated as such in their documentation.

a candidate entity used as a possession modifier, and its head in that relationship (the thing it possessed) was "city," then in comparing that with the children of a Named Entity Type graph's 'possession modifier' intermediary node branching from its root, we would find that its highest-scoring successor "town" gets a rating score of 0.77 vs. another successor "missiles," which gets a score of 0.11. So in this case we would add a score of 0.77 to this branch of this Named Entity Graph's similarity score, indicating that both of these entity's may have something city-like or town-like. We tried various alternative similarity scores, which we discuss below in Section 3.3, but found this most basic accumulation of similarity evidence to work the best in our practice so far.

4. We sum the scores of each "syntactic function node"'s highest-scoring child nodes to arrive at the similarity score for that Named Entity Type Graph.

5. We compare the score of each NET graph, and choose the highest-scoring type which best matched our candidate graph as our prediction. If we encounter a tie, we choose the Named Entity Type that appears most frequently in the corpus from among our top-scoring NETs.

3.3 Trying different similarity score approaches.

We tried a number of approaches of formulas for best evaluating the similarity of each syntactic branch of the graph. To show the equations, we will use the following conventions:

- L represents the NET graph's leaf node that corresponds syntactically with our candidate graph's leaf node which we are scoring, and has the highest similarity score of its sibling nodes to our candidate node.
- L_{count} represents the count of occurrences of that node
- Int_{count} represents the count of the intermediary node that is a predecessor to this node on the graph (so in the first generation of our graphs, Int represents the syntactic relationship of L to the entity, and Int_{count} represents how many times this entity type filled that syntactic role in our training corpus)

Initially, we tried to take into account how frequently L appeared in that syntactic relationship

to the head node. The intuitive motivation was that while a person might 'release' a prisoner, organizations probably more commonly 'release' press releases, and thus the frequency within each syntactic branch of each NET would be indicative of how commonly used that behavior is for that type of entity.

$$\log(L_{score}) + \log\left(\frac{L_{count}}{Int_{count}}\right)$$

This resulted in very poor performance (equivalent to random guessing), however, particularly on entity types such as 'PERSON' and 'ORG' which appear frequently, and therefore the matches were unduly penalized by the size of Int_{count} .

In order to relax the penalty of more-commonly occurring intermediary nodes, we tried this with the square root of both the count and the intermediary node count:

$$\log(L_{score}) + \log\left(\frac{L_{\sqrt{(count)}}}{Int_{\sqrt{(count)}}}\right)$$

We also tried a logarithmic approach there without the outer log probabilities, as our numbers weren't too close to zero as to make those essential:

$$L_{score} * \frac{L_{\log(count)}}{Int_{\log(count)}}$$

These methods improved performance on a small subset of candidates by not penalizing the more frequently occurring types as severely, but they were still penalized, and thus this method still resulted in worse performance than simply using the similarity counts on their own.

We note that by using similarity scores between word embeddings, we are potentially violating our principle of avoiding difficult-to-explain mathematical constructs. But we believe that the concept of a comparison measurement between the usage of words is explained easily enough to still be considered 'auditable', even if a lay person might not be able to generate those measurements arrived at by a cosine similarity.

4 Results and discussion

As we were focused on disambiguation rather than named entity chunking within this project, we measured our results on all entities where the named entity chunks we were provided by

our chunking system³ matched with the labeled chunks in the OntoNotes corpus.

Table 2 shows our results. Naturally, we were hoping for better performance from the model than we achieved, so we will take a moment to investigate where our model worked best, and where it tended to go wrong.

4.1 Where our model succeeded

We see that our model performed best on identifying 'PERSON' type entities according to the F1 score, and that 'PERSON' was among the top performers in terms of both precision and recall. Other top performers include 'ORG', 'GPE', 'DATE', and 'MONEY'.

In evaluating the metadata of our NET graphs, we see that all five of those NETs have among our highest average degree scores of over 2.2 in call cases (with all but 'MONEY' having an average degree of over 2.7). By contrast, our lowest performing NET types such as 'WORK_OF_ART', 'EVENT', and 'LAW' all have average degree scores of ~ 1.7 or below. We also see that our lowest performers each encountered fewer than 2,000 examples in our training set, and our top performers each encountered greater than 6,000 except for 'MONEY' which came in at a little over 3,000 nodes, but may have just been easier to classify.

We show how our performance improved with more nodes and higher average degree in Figure 2. This does suggest that our model was functional and became more accurate as it learned on more data, although improvements appear to begin tapering off in the addition of more nodes at around 5,000 nodes and in the increased average degree at around 2.4. We believe the additional features we discuss in Section 4.3 could allow for continued performance improvements as the size of our training data continued to scale.

4.2 Where our model failed

As the obvious inverse to where we succeeded, we note that our model failed most significantly where we lacked sufficient training examples. There was a significant difference in results with under 1,000 examples compared to those with 3,000 or more.

Additionally we realized early on in our exploration of this approach that most entity types are frequently personified, which creates more overlap in the verbs and adjectives to which they are

connected than we initially anticipated. While a person might 'run,' so too might an organization 'run a bake sale,' and a work of art might 'run up a high price at auction.' Therefore the verb of running is not so indicative of a 'PERSON' type as we initially imagined would be the case. In fact, it is difficult to conceive of a verb or adjective indicative of one class that could not be utilized by other entity types as well. Our model does have a difficult time discerning between a 'PERSON' and something that is 'PERSON'-ified.

Another limitation of our model is that it takes a significant amount of computation to generate the NET graphs and evaluate candidate documents. While we adjusted our architecture to be able to process jobs in parallel across a cluster in order to improve training time, it still took ≥ 18 hours for all jobs to finish across a Kubernetes cluster of 26 nodes with a combined 52 vCPUs and 195GB RAM. We expected this approach to not be the most speed-optimized technique given our graph-based approach rather than a vectorized approach taking advantage of matrix multiplication, but it made iteration and experimentation slower than anticipated.

Neural networks are notoriously data-hungry, but in comparison to our more 'brute force' approach, we see that trainable weights are in fact impressively efficient at generating high degrees of accuracy given the same amount of data. Given the stark differences in performance, we think future efforts to making Named Entity Disambiguation more auditable might be justifiably spent in making neural networks more interpretable rather than trying to circumvent those techniques with less flexible methods. Nevertheless, in the next section we present ideas on how our model might be improved.

For those interested, simplified representations of our NET Graphs generated on the first 500 documents in our training set are available at: <https://www.dropbox.com/sh/byarjdhore7lyn8/AAAVz1Yu9h39AVjQqBMfhrL3a?dl=0>

4.3 Possible enhancements

We believe that our model's performance could be significantly improved by implementation of one or more of the following methods:

- We acknowledge that one limitation of our model is the amount of surrounding signal

³spaCy's NER module

Entity Type	Precision	Recall	F1-Score	Support	Avg. Degree	No. Nodes
CARDINAL	23.45%	24.47%	23.95%	1590	2.59	5109
DATE	43.19%	27.18%	33.36%	2870	3.10	6078
EVENT	1.72%	8.62%	2.87%	116	1.74	1350
FAC	4.14%	12.08%	6.16%	149	1.75	1739
GPE	35.43%	23.91%	28.55%	3509	3.15	8976
LANGUAGE	1.68%	42.86%	3.23%	7	1.74	452
LAW	3.24%	9.23%	4.80%	65	1.59	701
LOC	3.65%	9.76%	5.32%	297	2.00	2312
MONEY	28.92%	45.13%	35.25%	822	2.23	3253
NORP	21.27%	28.49%	24.35%	1404	2.40	5293
ORDINAL	6.71%	22.01%	10.28%	359	2.01	1554
ORG	44.02%	19.50%	27.03%	3867	2.77	11135
PERCENT	20.91%	16.49%	18.44%	667	2.05	2171
PERSON	42.58%	36.80%	39.48%	2891	2.60	9640
PRODUCT	3.40%	9.85%	5.06%	132	1.84	1348
QUANTITY	3.82%	9.29%	5.41%	183	1.86	1471
TIME	5.30%	12.44%	7.43%	201	2.03	1478
WORK_OF_ART	1.67%	15.87%	3.02%	63	1.77	1743
Overall	34.04%	25.62%	28.05%	19192	2.18 (avg.)	3656 (avg.)

Table 2: Our results show that our overall performance was best on 'PERSON'-type entities. We also display some metadata of each entity's NET graph.

we are taking in as features. We think that further exploration of multiple generations of head and child words to incorporate, as well as the incorporation of the children of head nodes (which would allow us to take in direct objects when our entity is the subject of a sentence pointing directly to a root verb) would be valuable additional signal.

- We believe there might be benefits to reducing the number of nodes under each intermediary node (the nodes representing syntactic relationships), as long as we can do so without significantly reducing the signal that represents the variety of ways that a leaf node can follow from those nodes (especially seeing how important the average degree relates to performance). We hypothesize that by clustering leaf nodes on the Named Entity Template graphs, and having each cluster absorb some form of weight based on the number of nodes that falls into that cluster, we could dramatically speed up inference time, as well as provide a form of regularization. As we saw in our earlier example, the similarity of "city" and "town" is only 0.77 in comparison to an exact match of 1.0 with

city: a difference of nearly 30%. By finding a large number of cluster centers in the embedding space that no word would match *exactly* while more words would be *near*, we think we might be able to more gracefully generalize to new words without hitting such hot spots due to exact matches vs. near misses. If maintaining interpretability of the NET graphs as their own artifacts became too confusing with clusters, we could either find the word nearest the cluster center to represent it, or else lock our cluster centers to the most central of the existing leaf nodes via a Partitioning Around Medoids (k-medoids) algorithm.

- We posit that running a TF-IDF weighting of the leaf nodes would improve results by bringing out the most salient examples in each graph. Alternatively or additionally, we might prune stop words from the graphs as not containing distinguishing information.
- Another potential method we predict would improve the accuracy of our model is to replace certain types of strings with a common replacement. For example, we could replace

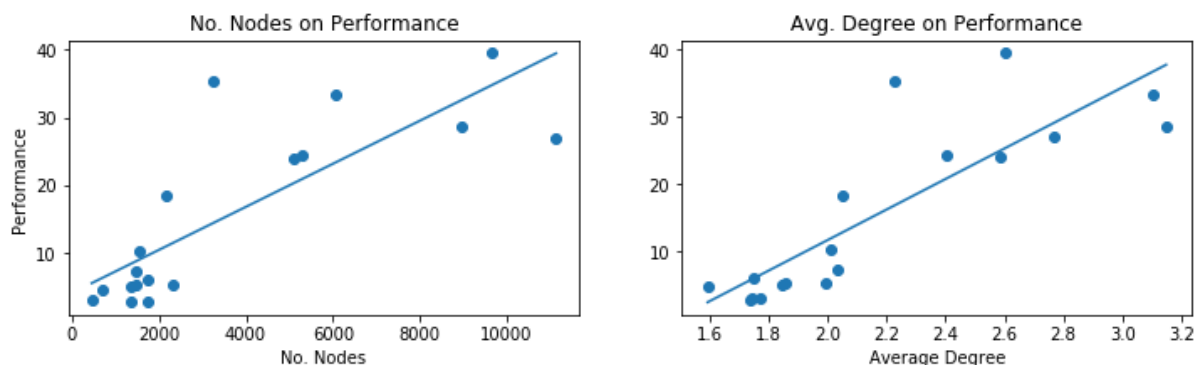


Figure 2: We see that performance improved both with more examples overall (no. of nodes) as well as greater average degree (specifically more examples per syntactic role).

all sequences of digits with a '1' in each digit location, so dates of all formats could result in '11/11/1111', prices could register as '\$1,111.11', and percentages could appear as '11%.' This replacement would be made in the NET graphs as well as in the candidate graphs. By reducing the large variance in numerical representations, our model would be more able to reliably interpret those types of data without the values needing to match exactly.

- A feature used by other NER systems that we believe might be valuable in ours is to incorporate the shape of input words, by which we mean a letter count that considers capitalization and digits vs text characters. For example 'Andrew' and 'Daniel' would both be represented as 'Xxxxxx', and '\$1,000' could be represented as '\$1,111'. By adding an additional 'shape' intermediary node as a parallel to the syntactic relationship intermediary nodes and perhaps giving it an additional (even trainable) weight, we expect we could gain additional performance, particularly on many of the numerical-based entity types of which we saw fewer examples.

5 Conclusion

Given more time to improve the inference performance of the model and to improve the features and potential weights included in the graph model, we believe that it could develop into a viable alternative for applications of NER where interpretability is critical.

Additionally, the graphs we generate are interesting artifacts in themselves, and we therefore see

possible uses for these techniques in evaluating in-progress literary works. One example use case might be in creating graphs that model the portrayal of each character in a novel, for example, to help publishers and authors clarify the distinct language choices distinguishing one character from another. With additional refinement and appropriately large labeled datasets, we could generate more fine-tuned distinctions as well, classifying a protagonist in a novel as a traditional hero or an anti-hero. We could identify occupations or literary tropes represented by supporting characters. Or we could graph primary characters on a per-chapter basis to assist in plotting the development of characters over a longer work.

Overall, we maintain a belief in the potential of graph-based 'mental constructs' as having additional applications in the progression of machine learning due to their persistence, general transferability from model to model, and their ability to continue to absorb cumulative knowledge. Nevertheless, we acknowledge the extreme performance advantages demonstrated by mature neural network approaches, and we believe that a more immediate (and potentially longterm) return on resources might come from developing tools to make those models more interpretable rather than attempting to circumvent the techniques that have led to their current prominence.

References

- Chris Dyer, Miguel Ballesteros, Wang Ling, Austin Matthews, and Noah A. Smith. 2015. Transition-Based Dependency Parsing with Stack Long Short-Term Memory.
- Matthew Honnibal and Ines Montani. 2017. spacy 2:

Natural language understanding with bloom embeddings, convolutional neural networks and incremental parsing. *To appear*.

Guillaume Lample, Miguel Ballesteros, Sandeep Subramanian, Kazuya Kawakami, and Chris Dyer. 2016. Neural Architectures for Named Entity Recognition.

Joakim Nivre. 2004. Incrementality in Deterministic Dependency Parsing. In *Proceedings of the Workshop on Incremental Parsing: Bringing Engineering and Cognition Together*, IncrementParsing '04, pages 50–57, Stroudsburg, PA, USA. Association for Computational Linguistics.

D. Suciú and A. Groza. 2014. [Interleaving ontology-based reasoning and Natural Language Processing for character identification in folktales](#). In *2014 IEEE 10th International Conference on Intelligent Computer Communication and Processing (ICCP)*, pages 67–74.